

CSC3150_Assignment3_Report 120090263

1.Environment of running my program

OS: CentOS Linux release 7.5.1804

VSC version: 1.73.0

CUDA version: 11.7

GPU information: Nvidia Quadro RTX 4000 GPU

Item	Configuration / Version
System Type	x86_64
Operating System	CentOS Linux release 7.5.1804
CPU	Intel(R) Xeon(R) Silver 4210R CPU @ 2.40GHz 20 Cores, 40 Threads
Memory	100GB RAM
GPU	Nvidia Quadro RTX 4000 GPU x 1
CUDA	11.7
GCC	Red Hat 7.3.1-5
CMake	3.14.1

2.Execution step (use VSCode)

1. enter "nvcc -rdc=true virtual_memory.cu main.cu user_program.cu -o virtual_memory"
2. enter "./virtual_memory"

```
[120090263@node21 After]$ nvcc -rdc=true virtual_memory.cu main.cu user_program.cu -o virtual_memory
```

```
[120090263@node21 After]$ ./virtual_memory
```

3.Design

In my program, I generally separate the situation into two cases.

The first case is page table is hit in read/write process. I write `check_hit()` function to go through page table, if the valid bit is valid and the virtual page number is matched, then the page is hit and return the frame number. In this case, if table is hit, I get the frame number and calculate *physical address*. Increase all frequency number by one and set the hit page to be 0. Then directly read/write the value.

read case1

```
//If page table hit, that is, the read content is in physical memory
frame_number = check_hit(vm, VPN);
empty_number = check_empty(vm);

//case1:if page table hit
if(frame_number != -1){
    phy_addr = REM + frame_number * vm->PAGESIZE;
    //将所有的frequency number加1, hit的设为0
    for(int i = 0; i < vm->PAGE_ENTRIES; i++){
        if(vm->invert_page_table[i] != 0x80000000){
            vm->invert_page_table[i + vm->PAGE_ENTRIES * 2] += 1;
        }
    }
    vm->invert_page_table[frame_number + vm->PAGE_ENTRIES * 2] = 0;
    ans = vm->buffer[phy_addr];

    return ans;
    // printf("hit times is %d\n", ++hit_time);
}
```

write case1

```
//Go through Page Table
frame_number = check_hit(vm, VPN);
empty_number = check_empty(vm);

//case1:if the page hit
if(frame_number != -1){
    phy_addr = REM + frame_number * vm->PAGESIZE;
    vm->buffer[phy_addr] = value;

    for(int i = 0; i < vm->PAGE_ENTRIES; i++){
        if(vm->invert_page_table[i] != 0x80000000){
            vm->invert_page_table[i + vm->PAGE_ENTRIES * 2] += 1;
        }
    }
    vm->invert_page_table[frame_number + vm->PAGE_ENTRIES * 2] = 0;
}
```

The second case is that page table is not hit. Further I separate it into two inner cases.

The first inner case is that the page table is not full (Although during read process this may not happen). In write function, simply update the page table, change frequency number and write the value into physical memory (virtual memory buffer).

The second inner case is that the page table is full. I use LRU to decide the entry to be swapped out of page table. That is, traverse the frequency number to get the largest one and its corresponding

frame number. First swap this page into disk and update the swap table. Then swap the requiring data from storage by locating the data through previously recorded swap table information.

3.1Details

3.1.1. Physical address calculation


To map virtual address to physical address, first obtain the frame number. I separate the given virtual address by dividing page size to get virtual page number and remaining byte. Then use the inverted page table to locate the VPN (If hit) and get the frame number. Then, frame number times page size and plus remaining byte to get the physical address.

```
u32 VPN = addr / vm->PAGESIZE;
u32 REM = addr % vm->PAGESIZE; //the remaining data in last page
```

3.1.2. LRU implementation

In this program, I implement the LRU using sequence number, that is, every entry in page table has an sequence number, as the process goes on, if the page is hit, the number will be set to 0, otherwise it increases. During the swap process, the one that has the largest frequency number should be swapped out. The frequency number is stored in page setting area.

```
__device__ void init_invert_page_table(VirtualMemory *vm) {
    for (int i = 0; i < vm->PAGE_ENTRIES; i++) {
        vm->invert_page_table[i] = 0x80000000; // invalid := MSB is 1
        vm->invert_page_table[i + vm->PAGE_ENTRIES] = i; //4-8KB store the VPN
        //initialize frequency number (8-12KB)
        vm->invert_page_table[i + vm->PAGE_ENTRIES * 2] = 0;
    }
}
```



```
//将所有的frequency number加1, hit的设为0
for(int i = 0; i < vm->PAGE_ENTRIES; i++){
    if(vm->invert_page_table[i] != 0x80000000){
        vm->invert_page_table[i + vm->PAGE_ENTRIES * 2] += 1;
    }
}
vm->invert_page_table[frame_number + vm->PAGE_ENTRIES * 2] = 0;
```

find the largest frequency number

```
u32 largest = vm->invert_page_table[vm->PAGE_ENTRIES * 2];

for(int i=0; i < vm->PAGE_ENTRIES; i++){//traverse to find the one that has largest frequency number and replace it
    if(vm->invert_page_table[i + vm->PAGE_ENTRIES * 2] >= largest){
        largest = vm->invert_page_table[i + vm->PAGE_ENTRIES * 2];
        frame_swap = i;
    }
}
//get the largest one
```

3.1.3. Swap Table implementation

Additionally, a swap table is designed. The virtual page number is used as index and the element is the address of swapped frame in disk. Everytime when one page is swapped out of disk, the empty pointer will point to the location. A boolean variable is used to deal with 2 situations that whether it is swap into disk or swap from disk to physical memory.

```
_device_ void Swap(VirtualMemory *vm, int frame_to_swap, u32 virtual_page_number, bool to_disk){
    if (to_disk == true){//use boolean to change whether swap into disk or from disk to pm
        //if swap into disk, update the storage index
        int tepVPN = vm->invert_page_table[frame_to_swap + vm->PAGE_ENTRIES];
        vm->swaptable[tepVPN] = (*vm->storage_index_ptr);
    }
    else{
        storage_swaptoPM = vm->swaptable[virtual_page_number]; //swap到 PM
        (*vm->storage_index_ptr) = storage_swaptoPM;
    }
}
```

swap process example in program

```
else{//inner case 2:---the page table is full
    (*vm->pagefault_num_ptr)++;
    // printf("This is VPN %d\n", VPN);
    u32 largest = 0;
    int frame_swap = 0;
    for(int i=0; i< vm->PAGE_ENTRIES; i++){//traverse to find the one that has largest frequency number and replace it
        if(vm->invert_page_table[i + vm->PAGE_ENTRIES * 2] >= largest){
            largest = vm->invert_page_table[i + vm->PAGE_ENTRIES * 2];
            frame_swap = i;
        }
    }
    //get the largest one

    Swap(vm, frame_swap, 0, true); // ←

    //SWAP largest frequency page back to disk-----
    for (int i = 0; i < vm->PAGESIZE; i++) {
        vm->storage[(*vm->storage_index_ptr)++] = vm->buffer[frame_swap * vm->PAGESIZE + i];
    }

    //SWAP the page from disk to physical memory

    Swap(vm, 0, VPN, false); // ←

    vm->invert_page_table[frame_swap + vm->PAGE_ENTRIES] = VPN; //update page table
    for(int i=0; i< vm->PAGESIZE; i++){
        vm->buffer[frame_swap * vm->PAGESIZE + i] = vm->storage[storage_swaptoPM + i];
    }

    phy_addr = REM + frame_swap * vm->PAGESIZE;
```

4. Page fault number

For test case 1, my program's page fault number is 8193. First, during write process, since there is nothing in physical memory, the page fault number is 4096. Then during read process, since it is reading from tail to head, the pages in physical memory will hit but as the data has one more byte

there will be one page fault. Last, the snapshot will again has 4096 page fault. Totally there will be 8193 pages.

For test case2, my program's page fault number is 9215. The first and second write process will have 4096 and 1023 page faults due to their data size respectively. Then the snapshot will have 4096 page faults. Totally there are 9215 page faults.

5. My problem

Problem: When designing the swap process, I encounter the problem that since I keep writing to disk regardless of those pages swapped back to physical memory, the disk will exceeds the given 128KB (131072 bytes).

Solution: Each time one page is swapped back to physical memory, I change the disk pointer to point at this location to indicate that this is empty. And next time when one page is swapped into disk, it can be stored in this location. Therefore, it wont exceed the given storage space.

6. Screen shot

test1

```
[120090263@node21 After]$ nvcc -rdc=true virtual_memory.cu main.cu user_program.cu -o virtual_memory
virtual_memory.cu(97): warning #177-D: variable "hit_time" was declared but never referenced

main.cu(97): warning #2464-D: conversion from a string literal to "char *" is deprecated
main.cu(116): warning #2464-D: conversion from a string literal to "char *" is deprecated
main.cu(97): warning #2464-D: conversion from a string literal to "char *" is deprecated
main.cu(116): warning #2464-D: conversion from a string literal to "char *" is deprecated

[120090263@node21 After]$ ./virtual_memory
input size: 131072
pagefault number is 8193
```

test2

```
[120090263@node21 After]$ nvcc -rdc=true virtual_memory.cu main.cu user_program.cu -o virtual_memory
virtual_memory.cu(97): warning #177-D: variable "hit_time" was declared but never referenced

main.cu(97): warning #2464-D: conversion from a string literal to "char *" is deprecated
main.cu(116): warning #2464-D: conversion from a string literal to "char *" is deprecated
main.cu(97): warning #2464-D: conversion from a string literal to "char *" is deprecated
main.cu(116): warning #2464-D: conversion from a string literal to "char *" is deprecated

[120090263@node21 After]$ ./virtual_memory
input size: 131072
pagefault number is 9215
```

7. What I learned from this project

In this project, I implement the virtual memory by using CUDA system. I get a better understanding of the relationship between virtual memory, physical memory and disk. I also understand how page table is implemented. Besides, I get to know the use of CUDA, including its structure, grammar and relation with C language.

***ps: Bonus task is not accomplished in this program**